

The Scheduling Algorithm

Eddie Guo

July 8, 2024

I created this scheduling algorithm to optimize call schedule for physicians and healthcare trainees. This system aims to create fair and efficient schedules while respecting various constraints and preferences.

Contents

1	Scheduling Constraints	1
1.1	No Back-to-Back Call Shifts	2
1.2	Minimum Interval Between Calls	2
1.3	Respecting Days Off	2
1.4	Fair Distribution of Shifts	2
1.5	Weekend Shift Optimization	2
1.6	Handling Constraint Conflicts	3
2	Main Function: generateCallSchedule	3
3	Key Components and Data Structures	3
3.1	DateTime Handling	3
3.2	Schedule Representation	3
3.3	Resident Data Structure	4
3.4	Tracking Variables	4
4	Scheduling Algorithm Overview	4
4.1	Resident Mode	4
4.2	Attending Mode	5
4.3	Schedule Optimization	5
5	Algorithm Output	5
6	Key Functions	5
6.1	canTakeCall	5
6.2	getResidentScore	6
7	Algorithm Details	6

1 Scheduling Constraints

The scheduling algorithm adheres to several important constraints to ensure fairness, compliance with labor regulations, and resident well-being. These constraints are implemented throughout the scheduling process and are crucial for generating a valid and effective schedule.

1.1 No Back-to-Back Call Shifts

One of the primary constraints is the prohibition of back-to-back call shifts for any resident. This rule is implemented as follows:

- When considering a resident for a call shift on a given date, the algorithm checks if the resident was assigned to a call shift on the previous day.
- If a resident was on call the previous day, they are automatically disqualified from being assigned to the current day's shift.
- This constraint is enforced in the `canTakeCall` function, which returns `false` if the resident was on call the previous day.

This constraint helps prevent physician fatigue and ensures compliance with work-hour regulations.

1.2 Minimum Interval Between Calls

Each resident has a `maxInX` property that defines the minimum number of days that must pass between their call shifts:

- The `canTakeCall` function checks the time elapsed since a resident's last call.
- If the time elapsed is less than the resident's `maxInX` value, the resident cannot be assigned to the current shift.
- This constraint allows for personalized scheduling based on resident seniority, specialization, or other factors that might influence their call frequency.

1.3 Respecting Days Off

The scheduler respects pre-defined days off for each resident:

- Each resident object includes a `daysOff` array containing dates when the resident is unavailable.
- The `canTakeCall` function checks if the current date is in the resident's `daysOff` array.
- If the date is found in `daysOff`, the resident is not considered for that shift.
- This constraint allows for accommodation of vacations, conferences, personal days, and other scheduled absences.

1.4 Fair Distribution of Shifts

While not a hard constraint, the scheduler aims to distribute shifts fairly among residents:

- The `getResidentScore` function considers the total number of shifts already assigned to a resident.
- Residents with fewer assigned shifts are given a higher score, increasing their chances of being selected for future shifts.
- This soft constraint helps balance the workload among residents over the entire scheduling period.

1.5 Weekend Shift Optimization

Depending on the `schedulingPreference`, the algorithm may treat weekend shifts differently:

- If the preference is set to `'minimize_weekend'`, the optimization phase focuses on reducing the number of weekend shifts for each resident.
- This is achieved by attempting to swap weekend assignments if it improves the overall schedule score.
- While not a hard constraint, this optimization aims to create a more equitable distribution of the often less desirable weekend shifts.

1.6 Handling Constraint Conflicts

In situations where it's impossible to satisfy all constraints (e.g., no available residents for a particular day), the scheduler employs a fallback strategy:

- It selects the least constrained resident, potentially violating the `max1inX` constraint if necessary.
- This ensures that all shifts are filled, even in challenging scheduling scenarios.
- The algorithm prioritizes filling all shifts over strict adherence to the `max1inX` constraint in these rare cases.

These constraints work together to create a schedule that is fair, respects resident preferences and limitations, complies with work regulations, and maintains high-quality patient care by ensuring well-rested physicians.

2 Main Function: `generateCallSchedule`

The primary function exported by this module is `generateCallSchedule`, which takes the following parameters:

- **`startDate`**: A JavaScript Date object representing the inclusive start date of the scheduling period.
- **`endDate`**: A JavaScript Date object representing the inclusive end date of the scheduling period.
- **`residents`**: An array of resident objects.
- **`schedulingPreference`**: A string indicating the preference for schedule optimization.
- **`schedulerMode`**: A string indicating whether to schedule for 'resident' or 'attending' mode.
- **`handoverDay`**: An integer (1-7) representing the day of the week for handovers in attending mode (default is 1 for Monday).

3 Key Components and Data Structures

3.1 DateTime Handling

The module leverages the `luxon` library for precise and consistent date and time operations. Specifically:

- **`DateTime` object**: Used to represent and manipulate dates throughout the scheduling process. This object provides methods for date arithmetic, formatting, and comparison, ensuring accurate handling of time-related operations.
- **Date Conversion**: JavaScript Date objects (used as input parameters) are converted to luxon `DateTime` objects for internal processing, then converted back to strings for the final schedule output.

3.2 Schedule Representation

The generated schedule is represented as a JavaScript object with the following structure:

- **Keys**: Strings representing dates in the 'YYYY-MM-DD' format. Each key corresponds to a single day in the scheduling period.
- **Values**: Strings representing resident names. Each value indicates the resident assigned to the corresponding date.

For example:

```
{
  '2024-07-01': 'Dr. Smith',
  '2024-07-02': 'Dr. Johnson',
  // ... more date-resident pairs
}
```

3.3 Resident Data Structure

Each resident in the `residents` array is represented by an object with the following properties:

- **name**: A string representing the resident's name. This is used as the identifier for the resident throughout the scheduling process.
- **maxInX**: An integer representing the minimum number of days that must pass between calls for this resident. For example, if `maxInX` is 4, the resident cannot be scheduled for a call more frequently than once every 4 days.
- **daysOff**: An array of date strings in the 'YYYY-MM-DD' format, representing days when the resident is unavailable for calls. This could include vacation days, conference attendance, or other pre-scheduled commitments.

Example of a resident object:

```
{
  name: 'Dr. Smith',
  maxInX: 4,
  daysOff: ['2024-07-04', '2024-07-05']
}
```

3.4 Tracking Variables

The algorithm maintains several important tracking variables throughout the scheduling process:

- **minIntervals**: An object that maps resident names to their respective `maxInX` values. This allows for quick lookup of a resident's minimum interval between calls.
- **lastCall**: An object that maps resident names to their most recently assigned call date. This is used to ensure the `maxInX` constraint is respected and to calculate scores for resident selection.
- **totalShifts**: An object that maps resident names to the total number of shifts they have been assigned so far. This is used to maintain fairness in shift distribution and for score calculation.

Examples of these tracking variables:

```
minIntervals = { 'Dr. Smith': 4, 'Dr. Johnson': 3 }
lastCall = { 'Dr. Smith': '2024-07-01', 'Dr. Johnson': '2024-06-30' }
totalShifts = { 'Dr. Smith': 5, 'Dr. Johnson': 6 }
```

4 Scheduling Algorithm Overview

The scheduling algorithm now has two distinct paths based on the `schedulerMode`:

4.1 Resident Mode

In resident mode, the algorithm fills shifts on a day-by-day basis:

1. Shuffle the `residents` array for fairness.
2. For each date in the scheduling period:
 - Filter available residents using `canTakeCall`.
 - If available residents exist, select the best candidate using `getResidentScore`.
 - Assign the selected resident to the current date.
 - Update tracking variables.

4.2 Attending Mode

In attending mode, the algorithm fills shifts on a weekly basis:

1. Determine the first and last full weeks based on the handover day.
2. Distribute full weeks attempting to exclude weeks with days off:
 - For each full week, find an available attending who can cover the entire week.
 - Assign the selected attending to all days in that week.
3. Handle partial weeks at the start and end separately:
 - Find available attendings for these partial periods.
 - Assign the attending with the least total shifts to cover the partial week.

4.3 Schedule Optimization

After the initial fill, this phase attempts to improve the schedule, focusing on weekend assignments:

1. Define weekend days based on the `schedulingPreference` parameter.
2. Iteratively attempt to improve weekend assignments:
 - For each weekend date, evaluate if swapping the assigned resident with another available resident would improve the overall score.
 - If an improvement is found, update the `schedule`, `lastCall`, and `totalShifts` accordingly.
 - Repeat this process until no further improvements can be made.

This optimization phase aims to create a more balanced and fair schedule, particularly for weekend shifts which are often considered less desirable.

5 Algorithm Output

The `generateCallSchedule` function returns a JavaScript object containing:

- `schedule`: The final generated schedule object, mapping dates to assigned residents.
- `totalShifts`: An object summarizing the total number of shifts assigned to each resident, providing a quick overview of shift distribution.

6 Key Functions

6.1 canTakeCall

Determines if a resident/attending can take a call on a given date:

$$\text{canTakeCall}(r, d, \text{duration}) = \begin{cases} \text{false} & \text{if } d \in r.\text{daysOff} \text{ or shift already assigned} \\ \text{false} & \text{if } d < \text{startDate} \text{ or } d > \text{endDate} \\ \text{true} & \text{if no previous call} \\ \text{interval} \geq r.\text{max1inX} & \text{if resident mode and } r.\text{max1inX} \text{ exists} \\ \text{interval} \geq 1 & \text{otherwise} \end{cases}$$

6.2 getResidentScore

Calculates a score for assigning a resident to a date:

$$\text{score} = \text{daysSinceLastCall} + 10 \left(\frac{\text{totalDays}}{\text{maxInX}} - \text{totalShifts} \right)$$

7 Algorithm Details

Algorithm 1 Generate Call Schedule

```
1: procedure GENERATECALLSCHEDULE(startDate, endDate, residents, schedulingPreference, scheduler-
   Mode, handoverDay)
2:   Initialize schedule, lastCall, totalShifts
3:   dates  $\leftarrow$  GenerateDateRange(startDate, endDate)
4:   ShuffleResidents(residents)
5:   if schedulerMode = 'resident' then
6:     FillResidentShifts(dates, residents, schedule)
7:   else if schedulerMode = 'attending' then
8:     FillAttendingShifts(dates, residents, schedule, handoverDay)
9:   end if
10:  warnings  $\leftarrow$  GenerateWarnings(schedule, schedulerMode)
11:  return schedule, totalShifts, warnings
12: end procedure
```

Algorithm 2 Fill Resident Shifts

```
1: procedure FILLRESIDENTSHIFTS(dates, residents, schedule)
2:   for each date in dates do
3:     availableResidents  $\leftarrow$  FilterAvailableResidents(residents, date)
4:     if availableResidents is not empty then
5:       selectedResident  $\leftarrow$  SelectBestResident(availableResidents, date)
6:       AssignShift(schedule, selectedResident, date)
7:     end if
8:   end for
9: end procedure
```

Algorithm 3 Fill Attending Shifts

```
1: procedure FILLATTENDINGSHIFTS(dates, residents, schedule, handoverDay)
2:   firstFullWeekStart, lastFullWeekEnd  $\leftarrow$  CalculateFullWeeks(dates, handoverDay)
3:   DistributeFullWeeks(firstFullWeekStart, lastFullWeekEnd, residents, schedule)
4:   HandlePartialWeekStart(dates[0], firstFullWeekStart, residents, schedule)
5:   HandlePartialWeekEnd(lastFullWeekEnd, dates[-1], residents, schedule)
6: end procedure
```
